

# THE PIGGYBANK

## DATA ENGINEERING CASE



The Piggybank is an American banking start-up. So far, bank has not that many customers, however it expects rapid growth in upcoming fiscal year. Currently all customers data is stored in CSV (Comma Separated Values) files and data is processed in the MS Office – all monthly reporting such as risk monitoring DPD30+, DPD60+, DPD90+ and overall credit portfolio summary.

The IT management has decided to switch from CSV files and use RDBMS (Relational Database Management System) to satisfy current and upcoming business needs.

You, as data engineer were asked to try out open-source RDBMS solution (SQLite), migrate all CSV files to it and recreate business reports that are running now in MS Office.

## Index

<b>I. Overview of used data and presentation of physical model.....</b>	<b>3</b>
a. Client – information about customers .....	3
b. Income – information about customer income .....	3
c. Household – information regarding customer family .....	3
d. Loan – information about loan status and characteristics .....	4
<b>II. Data quality checking .....</b>	<b>4</b>
a. CLIENT table.....	5
b. HOUSEHOLD table .....	6
c. INCOME table .....	7
d. LOAN table.....	8
<b>III. Entity relation diagram (ERD) for source tables .....</b>	<b>9</b>
<b>IV. Data Pipeline .....</b>	<b>10</b>
<b>V. RDBMS recommendation .....</b>	<b>11</b>

## I. Overview of used data and presentation of physical model

For the purpose of this business case, four data sets were provided and imported into a database.

### a. Client – information about customers

```
CREATE TABLE CLIENT (  
  CUSTOMER_ID INTEGER PRIMARY KEY AUTOINCREMENT,  
  REPORTING_DATE DATE NOT NULL,  
  AGE INTEGER (18, 100) NOT NULL,  
  EDUCATION VARCHAR (20) NOT NULL,  
  BUCKET INTEGER NOT NULL  
);
```

In the CLIENT table, CUSTOMER\_ID is used as a primary key since it allows to identify each row in the table with a unique value. An **index** is created for the column CUSTOMER\_ID to return customer information for a defined customer ID.

### b. Income – information about customer income

```
CREATE TABLE INCOME (  
  INCOME_ID INTEGER PRIMARY KEY AUTOINCREMENT,  
  CUSTOMER_ID INTEGER REFERENCES CLIENT (CUSTOMER_ID),  
  REPORTING_DATE DATE NOT NULL,  
  FIRST_JOB CHAR NOT NULL,  
  INCOME INTEGER (1000, 30000) NOT NULL,  
  BUCKET INTEGER NOT NULL  
);
```

In the INCOME table, INCOME\_ID is used as a primary key since it's a unique number attributed to customer income. A foreign key is set for CUSTOMER\_ID referring to CLIENT (CUSTOMER\_ID) to link both tables together based on the customer ID. An **index** is created for the columns CUSTOMER\_ID and INCOME\_ID to return customer income information based either on the customer or income ID.

### c. Household – information regarding customer family

```
CREATE TABLE HOUSEHOLD (  
  HOUSEHOLD_ID INTEGER PRIMARY KEY AUTOINCREMENT,  
  INCOME_ID INTEGER REFERENCES INCOME (INCOME_ID),  
  REPORTING_DATE DATE NOT NULL,  
  MARRIED CHAR NOT NULL,  
  HOUSE_OWNER CHAR NOT NULL,  
  CHILD_NO INTEGER (0, 10) NOT NULL,  
  HH_MEMBERS INTEGER (1, 10) NOT NULL,  
  BUCKET INTEGER NOT NULL  
);
```

In the HOUSEHOLD table, HOUSEHOLD\_ID is used as a primary key since the column represents a unique number that identifies the household. A foreign key is set for INCOME\_ID referring to INCOME (INCOME\_ID) to link both tables together based on the income ID. An **index** is created for the columns HOUSEHOLD\_ID and INCOME\_ID to return information regarding customer family based either on the household or income ID.

#### d. Loan – information about loan status and characteristics

```
CREATE TABLE LOAN (
  LOAN_ID          INTEGER          PRIMARY KEY AUTOINCREMENT,
  CUSTOMER_ID      INTEGER          REFERENCES CLIENT (CUSTOMER_ID),
  REPORTING_DATE   DATE             NOT NULL,
  INTODEFAULT      CHAR             NOT NULL,
  INSTALLMENT_NM   INTEGER (12, 72) NOT NULL,
  LOAN_AMT         DECIMAL (500, 100000) NOT NULL,
  INSTALLMENT_AMT  DECIMAL (10, 100000) NOT NULL,
  PAST_DUE_AMT     DECIMAL (0)      NOT NULL,
  BUCKET           INTEGER          NOT NULL
);
```

In the LOAN table, LOAN\_ID is used as a primary key since the column represents a unique number that identifies the loan. A foreign key is set for CUSTOMER\_ID referring to CLIENT (CUSTOMER\_ID) to link both tables together based on a customer's ID. An **index** is created for the columns LOAN\_ID and CUSTOMER\_ID to return information about loan characteristics based either on the loan or customer's ID.

## II. Data quality checking

In order to make accurate informed business decision, good quality of data is important. Poor data can lead to biased analysis and wrong recommendation. In this case study we will check data quality through 3 main metrics: **completeness**, **consistency**, and **accuracy**.

- Data completeness metric measures whether all the necessary data is present in the dataset. To do so, we can check the presence of null values in our tables.
- Data validity metric measures how well data conforms to required value attributes. To do so, we can check if the primary keys used are unique as they should.
- Data accuracy metric measures how accurately the data reflects reality. To do so, we can check if there are any outliers in the datasets. Inconsistencies in the data could occur during data entry and can lead to wrong results if not taken care of.

## a. CLIENT table

Checking the presence of null values in the table:

```
CREATE VIEW [client null assert] AS
SELECT *
FROM client
WHERE CUSTOMER_ID IS NULL OR
      REPORTING_DATE IS NULL OR
      AGE IS NULL OR
      EDUCATION IS NULL OR
      BUCKET IS NULL;
```


The query does not return any row meaning there is no null value in the table.

Verifying if the primary key (CUSTOMER\_ID field) is unique:

```
CREATE VIEW [client unique assert] AS
SELECT CUSTOMER_ID,
       SUM(1) AS count
FROM CLIENT
GROUP BY 1
HAVING count > 1;
```

The query does not return any row meaning there is no duplicate CUSTOMER\_ID values.

Trying to spot outliers in “AGE” column using plots:



The visual does not show any data inconsistency in the “AGE” column.

We then iterate to the other tables.

## b. HOUSEHOLD table

Checking the presence of null values in the table:

```
CREATE VIEW [household null assert] AS
SELECT *
FROM household
WHERE HOUSEHOLD_ID IS NULL OR
      INCOME_ID IS NULL OR
      REPORTING_DATE IS NULL OR
      MARRIED IS NULL OR
      HOUSE_OWNER IS NULL OR
      CHILD_NO IS NULL OR
      HH_MEMBERS IS NULL OR
      BUCKET IS NULL;
```

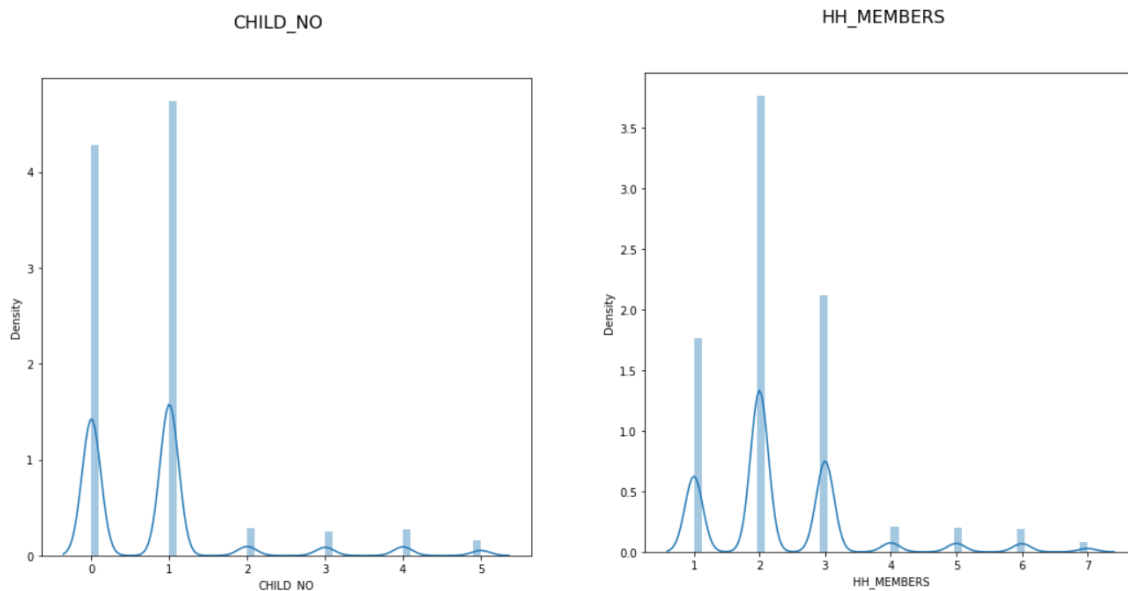
The query does not return any row meaning there is no null value in the table.

Verifying if the primary key (HOUSEHOLD\_ID field) is unique:

```
CREATE VIEW [household unique assert] AS
SELECT HOUSEHOLD_ID,
      SUM(1) AS count
FROM HOUSEHOLD
GROUP BY 1
HAVING count > 1;
```

The query does not return any row meaning there is no duplicate CUSTOMER\_ID values.

Trying to spot outliers in “CHILD\_NO” and “HH\_MEMBERS” columns using plots:



The visual does not show any data inconsistency in “CHILD\_NO” and “HH\_MEMBERS” columns.

### c. INCOME table

Checking the presence of null values in the table:

```
CREATE VIEW [income null assert] AS
SELECT *
FROM income
WHERE INCOME_ID IS NULL OR
      CUSTOMER_ID IS NULL OR
      REPORTING_DATE IS NULL OR
      FIRST_JOB IS NULL OR
      INCOME IS NULL OR
      BUCKET IS NULL;
```

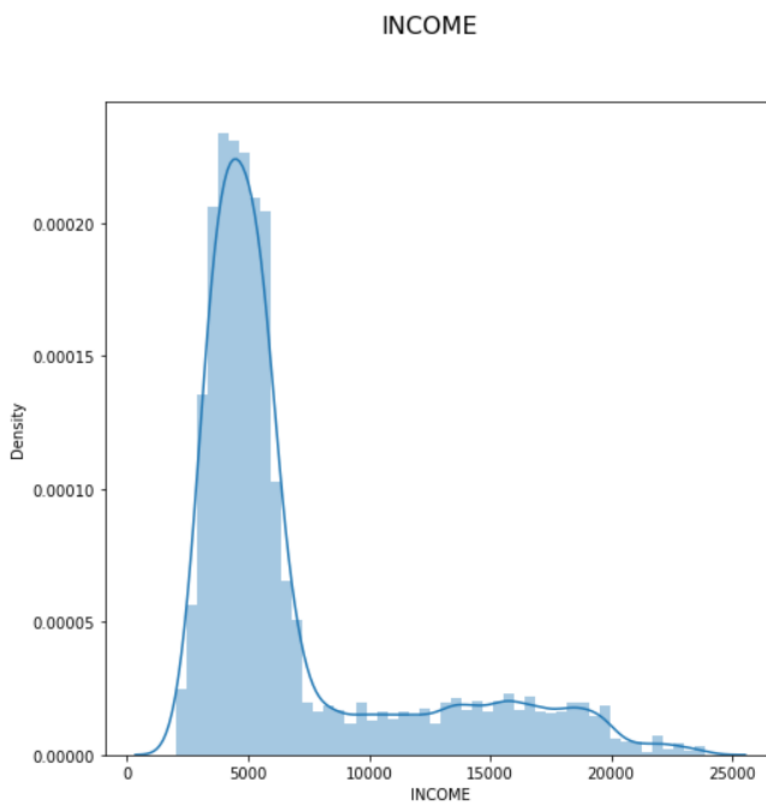
The query does not return any row meaning there is no null value in the table.

Verifying if the primary key (INCOME\_ID field) is unique:

```
CREATE VIEW [income unique assert] AS
SELECT INCOME_ID,
       SUM(1) AS count
FROM INCOME
GROUP BY 1
HAVING count > 1;
```

The query does not return any row meaning there is no duplicate INCOME\_ID values.

Trying to spot outliers in the “INCOME” column using plots:



The visual does not show any data inconsistency in the “INCOME\_NO” column.

#### d. LOAN table

Checking the presence of null values in the table:

```
CREATE VIEW [loan null assert] AS
SELECT *
FROM loan
WHERE LOAN_ID IS NULL OR
      CUSTOMER_ID IS NULL OR
      REPORTING_DATE IS NULL OR
      INTODEFAULT IS NULL OR
      INSTALLMENT_NM IS NULL OR
      LOAN_AMT IS NULL OR
      INSTALLMENT_AMT IS NULL OR
      PAST_DUE_AMT IS NULL OR
      BUCKET IS NULL;
```

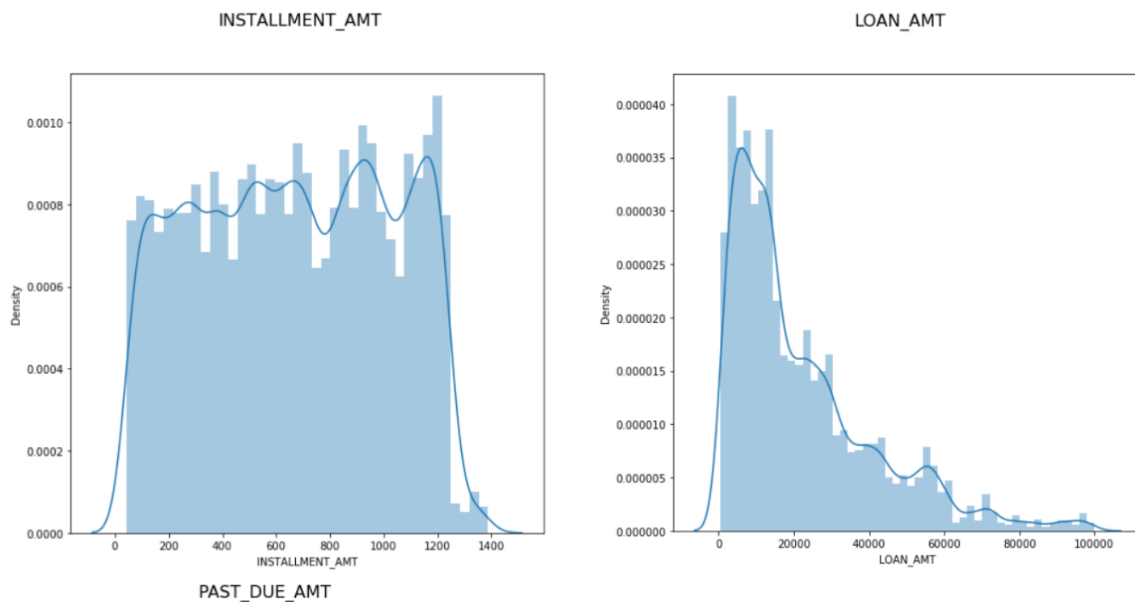
The query does not return any row meaning there is no null value in the table.

Verifying if the primary key (LOAN\_ID field) is unique:

```
CREATE VIEW [loan unique assert] AS
SELECT LOAN_ID,
       SUM(1) AS count
FROM LOAN
GROUP BY 1
HAVING count > 1;
```

The query does not return any row meaning there is no duplicate LOAN\_ID values.

Trying to spot outliers in “INSTALLMENT\_AMT” and “LOAN\_AMT” columns using plots:

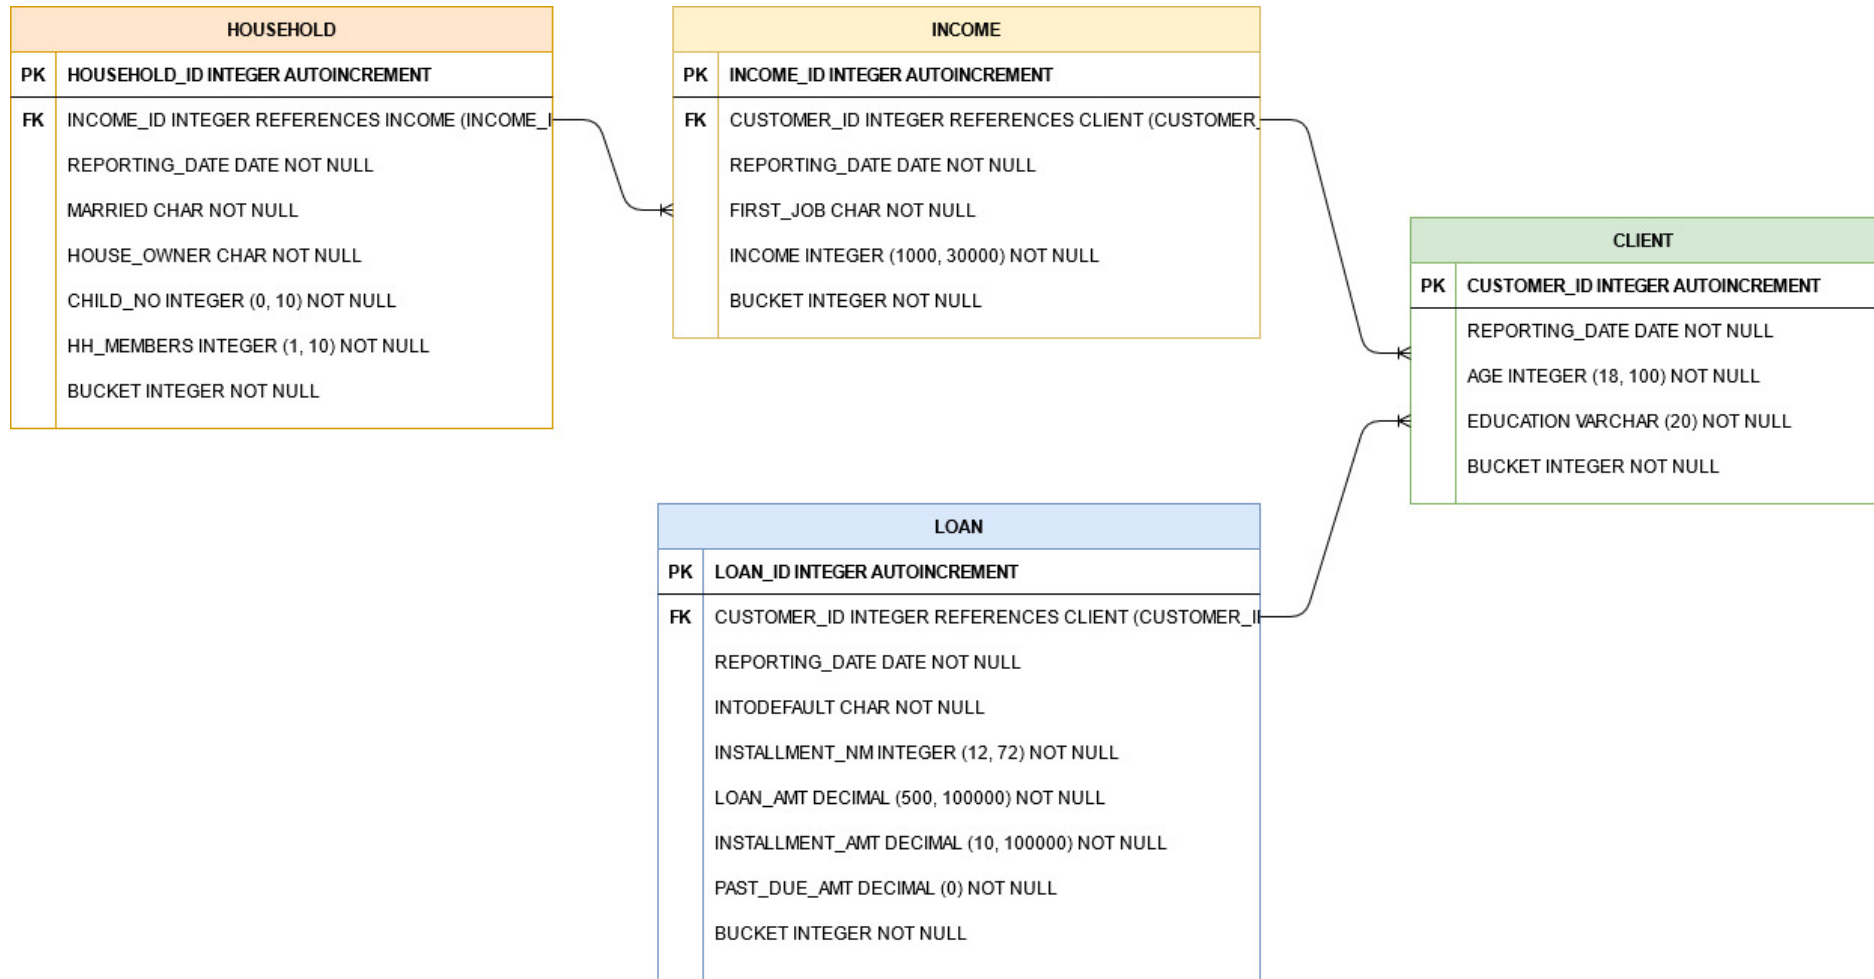


The visual does not show any data inconsistency in “INSTALLMENT\_AMT” and “LOAN\_AMT” columns.

The provided datasets do not require any cleaning since the data is complete, consistent, and accurate.



### III. Entity relation diagram (ERD) for source tables



## IV. Data Pipeline

The client expects an increasing amount of data for upcoming fiscal year. So, the pipeline must be scalable, ensure data quality and be able to collect and process data from different sources by matching and merging the data into a single record, the CUSTOMER\_CAR table (destination).

For a start-up in banking industry, we opt for a **Data Quality Pipeline** as it contains functions such as standardization of all new customer names at regular intervals and validation of customer's information like addresses that are useful for the loan application approval.

The data comes from the client's platform for loan application. As the clients are awaiting approval for their loan, and customer's data must be ready for analysis, the pipeline moves data from the source at regularly scheduled intervals of 20 minutes. The information collected is composed of data about the client, its household, its income and its loan situation.

The Pipeline performs standardization, matching and merging of the data to ensure the information matches the physical model of CUSTOMER\_CAR (format, schema).

In case of the need to handle data correction, it is possible to automate the correction in the "Transform" part of the ETL pipeline. For instance, if the date is entered in a different format, we can implement a replace tool. Same solution goes if a pattern is found in typos being present in the loaded data.

The clean and formatted data is then appended to CUSTOMER\_CAR.

## V. RDBMS recommendation

Since Piggybank is expecting a rapid growth of customer, the choice of RDBMS to store data is relevant. They are specifically designed to store and retrieve large amount of information, are consistent, stable, and reliable.

It is important to consider the business requirements before recommending a database system. The banking industry has its proper features and needs that must be taken into account. The ability to process daily volumes, to support high-rate transaction processing and data protection are obvious major concerns for a growing banking start-up.

In addition, Piggybank is looking for a RDBMS that satisfies other key requirements: the database must be open source, cost effective and offers flexibility, scalability, and reporting capabilities.

To fulfil those requirements, and based on the nature of the business, I would recommend two open-source RDBMS that are widely used in the banking sector: **MariaDB** and **PostgreSQL**.

PostgreSQL is a solution of choice for businesses of all sizes. By far the most mentioned advantage of PostgreSQL is the efficiency of its core algorithm, which means that it outperforms many databases supposed to be more advanced. This is especially useful if you are working with large datasets. It is also one of the most flexible open-source databases on the market. You can write functions in many languages: Python, Perl, Java, Ruby, C and R. Finally, as one of the most commonly used open-source databases, PostgreSQL's community support is one of the best on the market.

On the other hand, MariaDB is an open-source distribution of MySQL and was created after the acquisition of MySQL by Oracle. Many users choose MariaDB over MySQL because MariaDB has frequent security updates. While this does not necessarily mean that MariaDB is more secure, it does indicate that the development community takes security seriously, which represents an important point for a banking start-up. Other advantages of MariaDB are that it will almost certainly remain open source and highly compatible with MySQL. Because of this compatibility, MariaDB also works well with many other languages commonly used with MySQL.